

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
df=pd.read_csv(r"C:\Users\G S R KARTHIK\Downloads\fiat500_VehicleSelection_Dataset.csv")
print(df)
```

	ID	model	engine_power	age_in_days	km	previous_owners
0	1	lounge	51	882	25000	1 \
1	2	pop	51	1186	32500	1
2	3	sport	74	4658	142228	1
3	4	lounge	51	2739	160000	1
4	5	pop	73	3074	106880	1
...
1533	1534	sport	51	3712	115280	1
1534	1535	lounge	74	3835	112000	1
1535	1536	pop	51	2223	60457	1
1536	1537	lounge	51	2557	80750	1
1537	1538	pop	51	1766	54276	1

	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

[1538 rows x 9 columns]

In [2]:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing,svm
```

In [3]:

```
df=df[['km','price']]
df.columns=['Km','Price']
```

In [4]:



```
df.head(10)
```

Out[4]:

	Km	Price
0	25000	8900
1	32500	8800
2	142228	4200
3	160000	6000
4	106880	5700
5	70225	7900
6	11600	10750
7	49076	9190
8	76000	5600
9	89000	6000

In [5]:



```
df.tail()
```

Out[5]:

	Km	Price
1533	115280	5200
1534	112000	4600
1535	60457	7500
1536	80750	5990
1537	54276	7900

In [6]:



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Km      1538 non-null    int64  
 1   Price   1538 non-null    int64  
dtypes: int64(2)
memory usage: 24.2 KB
```

In [7]:



```
df.describe()
```

Out[7]:

	Km	Price
count	1538.000000	1538.000000
mean	53396.011704	8576.003901
std	40046.830723	1939.958641
min	1232.000000	2500.000000
25%	20006.250000	7122.500000
50%	39031.000000	9000.000000
75%	79667.750000	10000.000000
max	235000.000000	11100.000000

In [8]:



```
df.shape
```

Out[8]:

```
(1538, 2)
```

In [9]:



```
df.isnull().sum()
```

Out[9]:

```
Km      0
Price    0
dtype: int64
```

In [15]:



```
x=np.array(df['Km']).reshape(-1,1)
y=np.array(df['Price']).reshape(-1,1)
```

In [16]:



```
df.dropna(inplace=True)
```


In [17]:

```
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.7151460340273732

In [18]:

```
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```




In [20]:

```
udf=df[:][:500]  
sns.lmplot(x="Km",y="Price",data=udf,order=1,ci=None)
```

Out[20]:

<seaborn.axisgrid.FacetGrid at 0x1ec9b1279d0>



In [21]:


```
udf.fillna(method='ffill',inplace=True)  
X=np.array(udf['Km']).reshape(-1,1)  
y=np.array(udf['Price']).reshape(-1,1)  
udf.dropna(inplace=True)  
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3)  
regr.fit(X_train,y_train)
```

Out[21]:

```
LinearRegression  
LinearRegression()
```

In [22]:

```
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



In [24]:


```
from sklearn.linear_model import Ridge,Lasso,RidgeCV,LassoCV
```

In [25]:

```
plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

Out[25]:

<Axes: >



In [28]:

```
from sklearn.preprocessing import StandardScaler
features=df.columns[0:2]
target=df.columns[-1]
X=df[features].values
y=df[target].values
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
scaler=StandardScaler()
X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
```

The dimension of X_train is (1076, 2)
The dimension of X_test is (462, 2)

In [29]:

```
#Linear regression model
regr=LinearRegression()
regr.fit(X_train,y_train)
actual=y_test #actual value
train_score_regr=regr.score(X_train,y_train)
test_score_regr=regr.score(X_test,y_test)
print("\nLinear model:\n")
print("The train score for Linear model is {}".format(train_score_regr))
print("The test score for Linear model is {}".format(test_score_regr))
```

Linear model:

The train score for Linear model is 1.0
The test score for Linear model is 1.0

In [30]:

```
#ridge regression model
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge model:

The train score for ridge model is 0.9997095924476732
The test score for ridge model is 0.9997198323998524

In [31]:



```
#using the linear cv model for ridge regression
from sklearn.linear_model import RidgeCV
#ridge cross validation
ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
#score
print(ridge_cv.score(X_train,y_train))
print(ridge_cv.score(X_test,y_test))
```

0.9999999999999668

0.999999999999968

In [32]:




```
#using the linear cv model for lasso regression
from sklearn.linear_model import LassoCV
#lasso cross validation
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)
#score
print(lasso_cv.score(X_train,y_train))
print(lasso_cv.score(X_test,y_test))
```

0.9999999877496772

0.9999999874481674

In [33]:

```
plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,color='green')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```




In [35]:

```
#ridge regression  
plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

Out[35]:

<Axes: >



In [36]:



```
#Lasso regression model
lassoReg=Lasso(alpha=10)
lassoReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_lasso=lassoReg.score(X_train,y_train)
test_score_lasso=lassoReg.score(X_test,y_test)
print("\nLasso model:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```


Lasso model:

The train score for lasso model is 0.9999728562194999

The test score for lasso model is 0.9999728508562553

In [37]:

```
plt.figure(figsize=(10,10))
plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='o',markersize=5,color='red')
plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,color='green')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```




In [38]:



```
pd.Series(lassoReg.coef_, features).sort_values(ascending=True).plot(kind="bar")
```

Out[38]:

<Axes: >




In [39]:

```

#plot size
plt.figure(figsize=(10,10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
#add plot for lasso regression
plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='blue')
#add plot for linear model
plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize=7,color='green')
#rotate axis
plt.xticks(rotation=90)
plt.legend()
plt.title("Comparison of Ridge,Lasso and Linear regression models")
plt.show()

```



In []:

